

Import necessary libraries

```
import numpy as np
import pandas as pd
import seaborn as sns
import matplotlib.pyplot as plt
from sklearn.preprocessing import LabelEncoder
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler
from sklearn.linear_model import LogisticRegression
from sklearn.metrics import accuracy_score, classification_report,
confusion_matrix
```

Load the dataset

```
# Assuming the dataset is in a CSV file named 'Social_Network_Ads.csv'
df = pd.read_csv('Social_Network_Ads.csv')
```

Dataset Visualization

```
df.head()
```


	User ID	Gender	Age	EstimatedSalary	Purchased
0	15624510	Male	19	19000	0
1	15810944	Male	35	20000	0
2	15668575	Female	26	43000	0
3	15603246	Female	27	57000	0
4	15804002	Male	19	76000	0

```
df.shape
```


```
(400, 5)
```

EDA

```
ax = plt.subplots(figsize = (4,4))
ax = sns.countplot(x=df['Gender'])
plt.show()
```



```
ax = plt.subplots(figsize = (4,4))  
ax = sns.countplot(x=df['Purchased'])  
plt.show()
```



Feature Extraction

```
# Separate features (X) and target variable (y)
X = df.iloc[:, [1, 2, 3]].values # Considering Gender, Age, and
Estimated Salary as features
y = df.iloc[:, 4].values # Assuming 'Purchased' is the target
variable
```

Use LabelEncoder for 'Gender' as 'Gender' is non-numeric

```
label_encoder = LabelEncoder()
X[:, 0] = label_encoder.fit_transform(X[:, 0])
```

Split the dataset into training and testing sets

```
X_train, X_test, y_train, y_test = train_test_split(X, y,
test_size=0.2, random_state=42)
```

Feature scaling

```
scaler = StandardScaler()
X_train = scaler.fit_transform(X_train)
X_test = scaler.transform(X_test)
```

Create a Logistic Regression classifier

```
lr_classifier = LogisticRegression()
```

Fit the model to the training data

```
lr_classifier.fit(X_train, y_train)
LogisticRegression()
```

Make predictions on the test set

```
y_pred = lr_classifier.predict(X_test)
```

Evaluate the performance of the classifier

```
accuracy = accuracy_score(y_test, y_pred)
conf_matrix = confusion_matrix(y_test, y_pred)
classification_report_str = classification_report(y_test, y_pred)
```

Print the results

```
print(f'Logistic Regression Accuracy: {accuracy}')
print(f'Logistic Regression Confusion Matrix:\n{conf_matrix}')
sns.set(rc={'figure.figsize':(6,3)})
```

```
sns.heatmap(confusion_matrix(y_test,y_pred),annot = True,fmt = 'd')
plt.xlabel('Predicted Labels')
plt.ylabel('Actual Labels')
print(f'Logistic Regression Classification Report:\n{classification_report_str}')
```


Logistic Regression Accuracy: 0.8875

Logistic Regression Confusion Matrix:

```
[[50  2]
 [ 7 21]]
```

Logistic Regression Classification Report:

	precision	recall	f1-score	support
0	0.88	0.96	0.92	52
1	0.91	0.75	0.82	28
accuracy			0.89	80
macro avg	0.90	0.86	0.87	80
weighted avg	0.89	0.89	0.88	80



Predict whether a targeted audience or person will purchase the product or not

```
# Assuming we have a new set of feature values for prediction
new_data = np.array([[0, 30, 50000]]) # Example: Gender (0 for
Female, 1 for Male), Age, Estimated Salary

# Use the trained Logistic Regression model to make predictions
predicted_purchase = lr_classifier.predict(new_data)

# Print the predicted outcome
if predicted_purchase[0] == 1:
```

```
    print("The targeted audience is predicted to purchase the  
product.")  
else:  
    print("The targeted audience is predicted not to purchase the  
product.")
```

The targeted audience is predicted to purchase the product.

Output Visualization using Bar Plot

```
# Assuming we have already evaluated the model and obtained these  
metrics, hence plotting the same in a bar plot  
accuracy = 0.8875  
precision = 0.89  
recall = 0.89  
f1_score = 0.88  
  
# Plotting the bar plot  
metrics_names = ['Accuracy', 'Precision', 'Recall', 'F1-Score']  
metrics_values = [accuracy, precision, recall, f1_score]  
  
plt.bar(metrics_names, metrics_values, color=['blue', 'green',  
'orange', 'red'])  
plt.ylim([0, 1]) # Set the y-axis limit between 0 and 1  
plt.title('Model Evaluation Plot')  
plt.xlabel('Evaluation Metrics')  
plt.ylabel('Score')  
plt.show()
```



10-fold Cross-Validation

```

from sklearn.model_selection import cross_val_score, StratifiedKFold
from sklearn.metrics import make_scorer

# Define stratified 10-fold cross-validation
cross_val = StratifiedKFold(n_splits=10, shuffle=True,
random_state=42)

# Define accuracy as the evaluation metric
scoring = make_scorer(accuracy_score)

# Perform cross-validation on Logistic Regression
cv_results = cross_val_score(lr_classifier, X, y, cv=cross_val,
scoring=scoring)

# Display results
print("Cross-Validation Results:")
print("Individual Accuracies:", cv_results)
print("Average Accuracy:", np.mean(cv_results))

Cross-Validation Results:
Individual Accuracies: [0.65  0.65  0.65  0.65  0.65  0.65  0.65
0.625 0.625 0.625]
Average Accuracy: 0.6425000000000001

```

Cross-Validation Result Visualization using Bar Plot

```

# Cross-Validation Result
model = ['Logistic Regression']
accuracies = {
    'Logistic Regression': [0.65, 0.65, 0.65, 0.65, 0.65, 0.65, 0.65,
0.625, 0.625, 0.625],
}

# Plotting
plt.figure(figsize=(8, 4))

for model in model:
    plt.plot(range(1, 11), accuracies[model], marker='o',
label=f'{model} - Avg: {sum(accuracies[model])/10:.4f}')

plt.title('Individual Fold Accuracies of Logistic Regression
Classifier')
plt.xlabel('Trained Folds')
plt.ylabel('Accuracy')
plt.legend(bbox_to_anchor=(1.05, 1), loc='upper left') # Placing the
legend outside the plot area
plt.show()

```



ROC Curve Plotting for the above Logistic Regression Model

```
from sklearn.metrics import roc_curve, auc

# Assuming y_test is the actual labels
label_encoder = LabelEncoder()
y_test_binary = label_encoder.fit_transform(y_test)

# Get predicted probabilities for the positive class
lr_predicted_scores = lr_classifier.predict_proba(X_test)[: , 1]

# Compute ROC curve and AUC for Logistic Regression model
lr_fpr, lr_tpr, _ = roc_curve(y_test_binary, lr_predicted_scores)

# Compute AUC for Logistic Regression model
lr_roc_auc = auc(lr_fpr, lr_tpr)

# Plot ROC curve for Logistic Regression model
plt.figure(figsize=(8, 6))
sns.set(style='darkgrid')

plt.plot(lr_fpr, lr_tpr, color='purple', lw=2, label=f'Logistic Regression (AUC = {lr_roc_auc:.2f})')

plt.plot([0, 1], [0, 1], linestyle='--', color='gray', label='Random')
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('AUC-ROC Curve for Logistic Regression Model')
plt.legend(loc='lower right')
plt.show()
```

